

learn-co-curriculum / dsc-introduction-to-keras-lab Public

 View license

 1 star  176 forks

 Star

 Watch ▼

 Code  Issues  Pull requests  Actions  Projects  Security  Insights

 solution ▼

...

This branch is [11 commits ahead](#), [12 commits behind](#) master.



LoreDirick fixed env conflicts ...

on Feb 12, 2021

 12

[View code](#)

 README.md

Keras - Lab

Introduction

In this lab you'll once again build a neural network, but this time you will be using Keras to do a lot of the heavy lifting.

Objectives

You will be able to:

- Build a neural network using Keras
- Evaluate performance of a neural network using Keras

Required Packages

We'll start by importing all of the required packages and classes.

```
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import random
from sklearn.model_selection import train_test_split
from keras.utils.np_utils import to_categorical
from sklearn import preprocessing
from keras.preprocessing.text import Tokenizer
from keras import models
from keras import layers
from keras import optimizers
```

```
/Users/lore.dirick/anaconda3/lib/python3.6/site-packages/h5py/_init__.py:36:
FutureWarning: Conversion of the second argument of issubdtype from `float` to
`np.floating` is deprecated. In future, it will be treated as `np.float64 ==
np.dtype(float).type`.
    from ._conv import register_converters as _register_converters
Using TensorFlow backend.
```

Load the data

In this lab you will be classifying bank complaints available in the 'Bank_complaints.csv' file.

```
# Import data
df = pd.read_csv('Bank_complaints.csv')

# Inspect data
print(df.info())
df.head()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 60000 entries, 0 to 59999
Data columns (total 2 columns):
Product                  60000 non-null object
Consumer complaint narrative 60000 non-null object
dtypes: object(2)
memory usage: 937.6+ KB
None
```

```
<style scoped> .dataframe tbody tr th:only-of-type { vertical-align: middle; }
```

```
.dataframe tbody tr th {  
    vertical-align: top;  
}  
  
.dataframe thead th {  
    text-align: right;  
}  
  
</style>
```

	Product	Consumer complaint narrative
0	Student loan	In XX/XX/XXXX I filled out the Fedlaon applica...
1	Student loan	I am being contacted by a debt collector for p...
2	Student loan	I cosigned XXXX student loans at SallieMae for...
3	Student loan	Navient has sytematically and illegally failed...
4	Student loan	My wife became eligible for XXXX Loan Forgiven...

As mentioned earlier, your task is to categorize banking complaints into various predefined categories. Preview what these categories are and what percent of the complaints each accounts for.

```
df['Product'].value_counts(normalize=True)
```

```
Student loan          0.190067  
Credit card           0.159000  
Consumer Loan         0.157900  
Mortgage              0.138867  
Bank account or service 0.138483  
Credit reporting       0.114400  
Checking or savings account 0.101283  
Name: Product, dtype: float64
```

Preprocessing

Before we build our neural network, we need to do several preprocessing steps. First, we will create word vector counts (a bag of words type representation) of our complaints text. Next, we will change the category labels to integers. Finally, we will perform our usual train-test split before building and training our neural network using Keras. With that, let's start munging our data!

One-hot encoding of the complaints

Our first step again is to transform our textual data into a numerical representation. As we saw in some of our previous lessons on NLP, there are many ways to do this. Here, we'll use the `Tokenizer()` class from the `preprocessing.text` sub-module of the Keras package.

As with our previous work using NLTK, this will transform our text complaints into word vectors. (Note that the method of creating a vector is different from our previous work with NLTK; as you'll see, word order will be preserved as opposed to a bag of words representation). In the below code, we'll only keep the 2,000 most common words and use one-hot encoding.

```
# As a quick preliminary, briefly review the docstring for keras.preprocessing.text.  
Tokenizer?
```

```
# ⚠ This cell may take about thirty seconds to run
```

```
# Raw text complaints  
complaints = df['Consumer complaint narrative']  
  
# Initialize a tokenizer  
tokenizer = Tokenizer(num_words=2000)  
  
# Fit it to the complaints  
tokenizer.fit_on_texts(complaints)  
  
# Generate sequences  
sequences = tokenizer.texts_to_sequences(complaints)  
print('sequences type:', type(sequences))  
  
# Similar to sequences, but returns a numpy array  
one_hot_results = tokenizer.texts_to_matrix(complaints, mode='binary')  
print('one_hot_results type:', type(one_hot_results))
```

```
# Useful if we wish to decode (more explanation below)
word_index = tokenizer.word_index

# Tokens are the number of unique words across the corpus
print('Found %s unique tokens.' % len(word_index))

# Our coded data
print('Dimensions of our coded results:', np.shape(one_hot_results))
```

```
sequences type: <class 'list'>
one_hot_results type: <class 'numpy.ndarray'>
Found 50110 unique tokens.
Dimensions of our coded results: (60000, 2000)
```

Decoding Word Vectors

As a note, you can also decode these vectorized representations of the reviews. The `word_index` variable, defined above, stores the mapping from the label number to the actual word. Somewhat tediously, we can turn this dictionary inside out and map it back to our word vectors, giving us roughly the original complaint back. (As you'll see, the text won't be identical as we limited ourselves to top 2000 words.)

Python Review / Mini Challenge

While a bit tangential to our main topic of interest, we need to reverse our current dictionary `word_index` which maps words from our corpus to integers. In decoding our `one_hot_results`, we will need to create a dictionary of these integers to the original words. Below, take the `word_index` dictionary object and change the orientation so that the values are keys and the keys values. In other words, you are transforming something of the form {A:1, B:2, C:3} to {1:A, 2:B, 3:C}.

```
reverse_index = dict([(value, key) for (key, value) in word_index.items()])
```

Back to Decoding Our Word Vectors...

```
comment_idx_to_preview = 19
print('Original complaint text:')
print(complaints[comment_idx_to_preview])
```

```
print('\n\n')

#The reverse_index cell block above must be complete in order for this cell block to
decoded_review = ' '.join([reverse_index.get(i) for i in sequences[comment_idx_to_pr
print('Decoded review from Tokenizer:')
print(decoded_review)
```

Original complaint text:

I have already filed several complaints about AES/PHEAA. I was notified by a XXXX XXXX let @ XXXX, who pretended to be from your office, he said he was from CFPB. I found out this morning he is n't from your office, but is actually works at XXXX.

This has wasted weeks of my time. They AES/PHEAA confirmed and admitted (see attached transcript of XXXX, conversation at XXXX (XXXX) with XXXX that proves they verified the loans are not mine) the student loans they had XXXX, and collected on, and reported negate credit reporting in my name are in fact, not mine.

They conclued their investigation on XXXX admitting they made a mistake and have my name on soneone elses loans. I these XXXX loans total {\$10000.00}, original amount. My XXXX loans I got was total {\$3500.00}. We proved by providing AES/PHEAA, this with my original promissary notes I located recently, the XXXX of my college provided AES/PHEAA with their original shoeinf amounts of my XXXX loans which show different dates and amounts, the dates and amounts are not even close to matching these loans they have in my name, The original lender, XXXX XXXX Bank notifying AES/PHEAA, they never issued me a student loan, and original Loan Guarantor, XXXX, notifying AES/PHEAA, they never were guarantor of my loans.

XXXX straight forward. But today, this person, XXXX XXXX, told me they know these loans are not mine, and they refuse to remove my name off these XXXX loan 's and correct their mistake, essentially forcing me to pay these loans off, bucause in XXXX they sold the loans to XXXX loans.

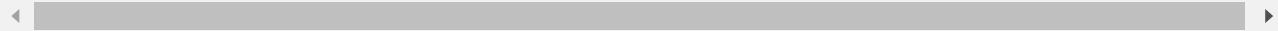
This is absurd, first protruding to be this office, and then refusing to correct their mistake.

Please for the love of XXXX will soneone from your office call me at XXXX, today. I am a XXXX vet and they are knowingly discriminating against me. Pretending to be you.

Decoded review from Tokenizer:

i have already filed several complaints about aes i was notified by a xxxx xxxx let xxxx who to be from your office he said he was from cfpb i found out this morning he is n't from your office but is actually works at xxxx this has weeks of my time they aes confirmed and admitted see attached of xxxx conversation at

xxxx xxxx with xxxx that they verified the loans are not mine the student loans they had xxxx and on and reported credit reporting in my name are in fact not mine they their investigation on xxxx they made a mistake and have my name on loans i these xxxx loans total 10000 00 original amount my xxxx loans i got was total 00 we by providing aes this with my original notes i located recently the xxxx of my college provided aes with their original amounts of my xxxx loans which show different dates and amounts the dates and amounts are not even close to these loans they have in my name the original lender xxxx xxxx bank notifying aes they never issued me a student loan and original loan xxxx notifying aes they never were of my loans xxxx forward but today this person xxxx xxxx told me they know these loans are not mine and they refuse to remove my name off these xxxx loan 's and correct their mistake essentially me to pay these loans off in xxxx they sold the loans to xxxx loans this is first to be this office and then refusing to correct their mistake please for the of xxxx will from your office call me at xxxx today i am a xxxx and they are against me to be you



Convert the Products to Numerical Categories

On to step two of our preprocessing: converting our descriptive categories into integers.

```
product = df['Product']

# Initialize
le = preprocessing.LabelEncoder()
le.fit(product)
print('Original class labels:')
print(list(le.classes_))
print('\n')
product_cat = le.transform(product)

# If you wish to retrieve the original descriptive labels post production
# list(le.inverse_transform([0, 1, 3, 3, 0, 6, 4]))

print('New product labels:')
print(product_cat)
print('\n')

# Each row will be all zeros except for the category for that observation
print('One hot labels; 7 binary columns, one for each of the categories.')
product_onehot = to_categorical(product_cat)
print(product_onehot)
print('\n')
```

```
print('One hot labels shape:')
```

```
print(np.shape(product_onehot))
```



```
Original class labels:
```

```
['Bank account or service', 'Checking or savings account', 'Consumer Loan',
```

```
'Credit card', 'Credit reporting', 'Mortgage', 'Student loan']
```

```
New product labels:
```

```
[6 6 6 ... 4 4 4]
```

```
One hot labels; 7 binary columns, one for each of the categories.
```

```
[[0. 0. 0. ... 0. 0. 1.]
```

```
[0. 0. 0. ... 0. 0. 1.]
```

```
[0. 0. 0. ... 0. 0. 1.]
```



```
...
```

```
[0. 0. 0. ... 1. 0. 0.]
```

```
[0. 0. 0. ... 1. 0. 0.]
```

```
[0. 0. 0. ... 1. 0. 0.]]
```

```
One hot labels shape:
```

```
(60000, 7)
```

Train-test split

Now for our final preprocessing step: the usual train-test split.

```
random.seed(123)
test_index = random.sample(range(1,10000), 1500)

test = one_hot_results[test_index]
train = np.delete(one_hot_results, test_index, 0)

label_test = product_onehot[test_index]
label_train = np.delete(product_onehot, test_index, 0)

print('Test label shape:', np.shape(label_test))
print('Train label shape:', np.shape(label_train))
print('Test shape:', np.shape(test))
print('Train shape:', np.shape(train))
```

```
Test label shape: (1500, 7)
Train label shape: (58500, 7)
Test shape: (1500, 2000)
Train shape: (58500, 2000)
```

Building the network

Let's build a fully connected (Dense) layer network with relu activation in Keras. You can do this using: `Dense(16, activation='relu')`.

In this example, use two hidden layers with 50 units in the first layer and 25 in the second, both with a 'relu' activation function. Because we are dealing with a multiclass problem (classifying the complaints into 7 categories), we use a use a 'softmax' classifier in order to output 7 class probabilities per case.

```
# Initialize a sequential model
model = models.Sequential()

# Two layers with relu activation
model.add(layers.Dense(50, activation='relu', input_shape=(2000,)))
model.add(layers.Dense(25, activation='relu'))

# One layer with softmax activation
model.add(layers.Dense(7, activation='softmax'))
```

Compiling the model

Now, compile the model! This time, use 'categorical_crossentropy' as the loss function and stochastic gradient descent, 'SGD' as the optimizer. As in the previous lesson, include the accuracy as a metric.

```
# Compile the model
model.compile(optimizer='SGD',
              loss='categorical_crossentropy',
              metrics=['acc'])
```

Training the model

In the compiler, you'll be passing the optimizer (SGD = stochastic gradient descent), loss function, and metrics. Train the model for 120 epochs in mini-batches of 256 samples.

Note:  Your code may take about one to two minutes to run.

```
# Train the model
history = model.fit(train,
                     label_train,
                     epochs=120,
                     batch_size=256)
```

```
Epoch 1/120
58500/58500 [=====] - 2s 40us/step - loss: 1.8825 - acc: 0.2408
Epoch 2/120
58500/58500 [=====] - 2s 36us/step - loss: 1.6170 - acc: 0.4628
Epoch 3/120
58500/58500 [=====] - 2s 31us/step - loss: 1.2554 - acc: 0.6184
Epoch 4/120
58500/58500 [=====] - 2s 32us/step - loss: 0.9964 - acc: 0.6802
Epoch 5/120
58500/58500 [=====] - 2s 28us/step - loss: 0.8516 - acc: 0.7086
Epoch 6/120
58500/58500 [=====] - 2s 29us/step - loss: 0.7686 - acc: 0.7263
Epoch 7/120
58500/58500 [=====] - 2s 28us/step - loss: 0.7161 - acc: 0.7390
Epoch 8/120
58500/58500 [=====] - 2s 33us/step - loss: 0.6796 - acc: 0.7493
Epoch 9/120
58500/58500 [=====] - 2s 32us/step - loss: 0.6521 - acc: 0.7575
Epoch 10/120
58500/58500 [=====] - 2s 29us/step - loss: 0.6302 - acc: 0.7662
Epoch 11/120
58500/58500 [=====] - 2s 30us/step - loss: 0.6125 - acc: 0.7724
```

```
Epoch 12/120
58500/58500 [=====] - 1s 22us/step - loss: 0.5973 - acc: 0.7782
Epoch 13/120
58500/58500 [=====] - 1s 23us/step - loss: 0.5844 - acc: 0.7827
Epoch 14/120
58500/58500 [=====] - 1s 25us/step - loss: 0.5728 - acc: 0.7871
Epoch 15/120
58500/58500 [=====] - 1s 25us/step - loss: 0.5622 - acc: 0.7917
Epoch 16/120
58500/58500 [=====] - 2s 28us/step - loss: 0.5527 - acc: 0.7965: 0s - loss: 0.5514 - acc:
Epoch 17/120
58500/58500 [=====] - 1s 23us/step - loss: 0.5440 - acc: 0.7994
Epoch 18/120
58500/58500 [=====] - 1s 23us/step - loss: 0.5359 - acc: 0.8037
Epoch 19/120
58500/58500 [=====] - 1s 22us/step - loss: 0.5286 - acc: 0.8072
Epoch 20/120
58500/58500 [=====] - 1s 21us/step - loss: 0.5215 - acc: 0.8098
Epoch 21/120
58500/58500 [=====] - 1s 23us/step - loss: 0.5152 - acc: 0.8128
Epoch 22/120
58500/58500 [=====] - 1s 25us/step - loss: 0.5092 - acc: 0.8148
Epoch 23/120
58500/58500 [=====] - 1s 22us/step - loss: 0.5036 - acc: 0.8170
Epoch 24/120
58500/58500 [=====] - 1s 24us/step - loss: 0.4977 - acc: 0.8198
Epoch 25/120
58500/58500 [=====] - 1s 23us/step - loss: 0.4930 - acc: 0.8224
Epoch 26/120
58500/58500 [=====] - 1s 21us/step - loss: 0.4882 - acc: 0.8238
Epoch 27/120
58500/58500 [=====] - 1s 19us/step - loss: 0.4836 - acc: 0.8256
Epoch 28/120
```

```
58500/58500 [=====] - 1s 20us/step - loss: 0.4793 - acc: 0.8275
Epoch 29/120
58500/58500 [=====] - 1s 20us/step - loss: 0.4752 - acc: 0.8296
Epoch 30/120
58500/58500 [=====] - 1s 21us/step - loss: 0.4712 - acc: 0.8305
Epoch 31/120
58500/58500 [=====] - 1s 22us/step - loss: 0.4676 - acc: 0.8319
Epoch 32/120
58500/58500 [=====] - 1s 21us/step - loss: 0.4638 - acc: 0.8342
Epoch 33/120
58500/58500 [=====] - 1s 25us/step - loss: 0.4607 - acc: 0.8351
Epoch 34/120
58500/58500 [=====] - 1s 24us/step - loss: 0.4577 - acc: 0.8362
Epoch 35/120
58500/58500 [=====] - 2s 28us/step - loss: 0.4543 - acc: 0.8380
Epoch 36/120
58500/58500 [=====] - 1s 20us/step - loss: 0.4511 - acc: 0.8389
Epoch 37/120
58500/58500 [=====] - 1s 20us/step - loss: 0.4483 - acc: 0.8397
Epoch 38/120
58500/58500 [=====] - 1s 20us/step - loss: 0.4454 - acc: 0.8407
Epoch 39/120
58500/58500 [=====] - 1s 20us/step - loss: 0.4425 - acc: 0.8426
Epoch 40/120
58500/58500 [=====] - 1s 19us/step - loss: 0.4401 - acc: 0.8434
Epoch 41/120
58500/58500 [=====] - 1s 17us/step - loss: 0.4375 - acc: 0.8435
Epoch 42/120
58500/58500 [=====] - 1s 20us/step - loss: 0.4349 - acc: 0.8449
Epoch 43/120
58500/58500 [=====] - 1s 20us/step - loss: 0.4328 - acc: 0.8457
Epoch 44/120
58500/58500 [=====] - 1s 17us/step - loss: 0.4303 - acc:
```

```
0.8469
Epoch 45/120
58500/58500 [=====] - 1s 19us/step - loss: 0.4283 - acc:
0.8480
Epoch 46/120
58500/58500 [=====] - 1s 19us/step - loss: 0.4262 - acc:
0.8475
Epoch 47/120
58500/58500 [=====] - 1s 17us/step - loss: 0.4237 - acc:
0.8490
Epoch 48/120
58500/58500 [=====] - 1s 20us/step - loss: 0.4221 - acc:
0.8495
Epoch 49/120
58500/58500 [=====] - 1s 18us/step - loss: 0.4201 - acc:
0.8508
Epoch 50/120
58500/58500 [=====] - 1s 16us/step - loss: 0.4182 - acc:
0.8508
Epoch 51/120
58500/58500 [=====] - 1s 15us/step - loss: 0.4163 - acc:
0.8511
Epoch 52/120
58500/58500 [=====] - 1s 18us/step - loss: 0.4140 - acc:
0.8518
Epoch 53/120
58500/58500 [=====] - 1s 19us/step - loss: 0.4125 - acc:
0.8532
Epoch 54/120
58500/58500 [=====] - 1s 19us/step - loss: 0.4106 - acc:
0.8531
Epoch 55/120
58500/58500 [=====] - 1s 18us/step - loss: 0.4091 - acc:
0.8544
Epoch 56/120
58500/58500 [=====] - 1s 18us/step - loss: 0.4077 - acc:
0.8546
Epoch 57/120
58500/58500 [=====] - 1s 19us/step - loss: 0.4057 - acc:
0.8550
Epoch 58/120
58500/58500 [=====] - 1s 19us/step - loss: 0.4040 - acc:
0.8558
Epoch 59/120
58500/58500 [=====] - 1s 19us/step - loss: 0.4029 - acc:
0.8564
Epoch 60/120
58500/58500 [=====] - 1s 21us/step - loss: 0.4009 - acc:
0.8565
```

```
Epoch 61/120
58500/58500 [=====] - 1s 23us/step - loss: 0.3998 - acc: 0.8577
Epoch 62/120
58500/58500 [=====] - 1s 22us/step - loss: 0.3984 - acc: 0.8578
Epoch 63/120
58500/58500 [=====] - 1s 17us/step - loss: 0.3966 - acc: 0.8584
Epoch 64/120
58500/58500 [=====] - 1s 19us/step - loss: 0.3953 - acc: 0.8593
Epoch 65/120
58500/58500 [=====] - 1s 20us/step - loss: 0.3939 - acc: 0.8591
Epoch 66/120
58500/58500 [=====] - 1s 19us/step - loss: 0.3931 - acc: 0.8602
Epoch 67/120
58500/58500 [=====] - 1s 19us/step - loss: 0.3913 - acc: 0.8602
Epoch 68/120
58500/58500 [=====] - 1s 19us/step - loss: 0.3900 - acc: 0.8608
Epoch 69/120
58500/58500 [=====] - 1s 20us/step - loss: 0.3889 - acc: 0.8614
Epoch 70/120
58500/58500 [=====] - 1s 21us/step - loss: 0.3873 - acc: 0.8613
Epoch 71/120
58500/58500 [=====] - 1s 21us/step - loss: 0.3861 - acc: 0.8622
Epoch 72/120
58500/58500 [=====] - 1s 21us/step - loss: 0.3851 - acc: 0.8624
Epoch 73/120
58500/58500 [=====] - 1s 19us/step - loss: 0.3838 - acc: 0.8630
Epoch 74/120
58500/58500 [=====] - 1s 20us/step - loss: 0.3826 - acc: 0.8631
Epoch 75/120
58500/58500 [=====] - 1s 19us/step - loss: 0.3815 - acc: 0.8640
Epoch 76/120
58500/58500 [=====] - 1s 20us/step - loss: 0.3806 - acc: 0.8642
Epoch 77/120
```

```
58500/58500 [=====] - 1s 19us/step - loss: 0.3792 - acc: 0.8640
Epoch 78/120
58500/58500 [=====] - 1s 20us/step - loss: 0.3780 - acc: 0.8648
Epoch 79/120
58500/58500 [=====] - 1s 21us/step - loss: 0.3770 - acc: 0.8650
Epoch 80/120
58500/58500 [=====] - 1s 21us/step - loss: 0.3758 - acc: 0.8655
Epoch 81/120
58500/58500 [=====] - 1s 20us/step - loss: 0.3750 - acc: 0.8650
Epoch 82/120
58500/58500 [=====] - 1s 19us/step - loss: 0.3742 - acc: 0.8665
Epoch 83/120
58500/58500 [=====] - 1s 24us/step - loss: 0.3727 - acc: 0.8660
Epoch 84/120
58500/58500 [=====] - 1s 19us/step - loss: 0.3720 - acc: 0.8671
Epoch 85/120
58500/58500 [=====] - 1s 18us/step - loss: 0.3710 - acc: 0.8673
Epoch 86/120
58500/58500 [=====] - 1s 19us/step - loss: 0.3700 - acc: 0.8667
Epoch 87/120
58500/58500 [=====] - 1s 20us/step - loss: 0.3690 - acc: 0.8682
Epoch 88/120
58500/58500 [=====] - 1s 21us/step - loss: 0.3680 - acc: 0.8685
Epoch 89/120
58500/58500 [=====] - 1s 19us/step - loss: 0.3673 - acc: 0.8693
Epoch 90/120
58500/58500 [=====] - 1s 25us/step - loss: 0.3660 - acc: 0.8686
Epoch 91/120
58500/58500 [=====] - 1s 24us/step - loss: 0.3653 - acc: 0.8693
Epoch 92/120
58500/58500 [=====] - 1s 21us/step - loss: 0.3643 - acc: 0.8695
Epoch 93/120
58500/58500 [=====] - 1s 20us/step - loss: 0.3635 - acc:
```

```
0.8705
Epoch 94/120
58500/58500 [=====] - 1s 21us/step - loss: 0.3623 - acc:
0.8706
Epoch 95/120
58500/58500 [=====] - 1s 21us/step - loss: 0.3617 - acc:
0.8705
Epoch 96/120
58500/58500 [=====] - 1s 20us/step - loss: 0.3608 - acc:
0.8710
Epoch 97/120
58500/58500 [=====] - 1s 20us/step - loss: 0.3600 - acc:
0.8715
Epoch 98/120
58500/58500 [=====] - 1s 21us/step - loss: 0.3590 - acc:
0.8711
Epoch 99/120
58500/58500 [=====] - 1s 23us/step - loss: 0.3579 - acc:
0.8724
Epoch 100/120
58500/58500 [=====] - 1s 20us/step - loss: 0.3573 - acc:
0.8725
Epoch 101/120
58500/58500 [=====] - 1s 19us/step - loss: 0.3563 - acc:
0.8722
Epoch 102/120
58500/58500 [=====] - 1s 20us/step - loss: 0.3555 - acc:
0.8726
Epoch 103/120
58500/58500 [=====] - 1s 20us/step - loss: 0.3549 - acc:
0.8729
Epoch 104/120
58500/58500 [=====] - 1s 20us/step - loss: 0.3539 - acc:
0.8734
Epoch 105/120
58500/58500 [=====] - 1s 19us/step - loss: 0.3531 - acc:
0.8743
Epoch 106/120
58500/58500 [=====] - 1s 19us/step - loss: 0.3522 - acc:
0.8745
Epoch 107/120
58500/58500 [=====] - 1s 20us/step - loss: 0.3514 - acc:
0.8746
Epoch 108/120
58500/58500 [=====] - 1s 19us/step - loss: 0.3507 - acc:
0.8750
Epoch 109/120
58500/58500 [=====] - 1s 20us/step - loss: 0.3499 - acc:
0.8747
```

```
Epoch 110/120
58500/58500 [=====] - 1s 19us/step - loss: 0.3487 - acc: 0.8750
Epoch 111/120
58500/58500 [=====] - 1s 20us/step - loss: 0.3484 - acc: 0.8752
Epoch 112/120
58500/58500 [=====] - 1s 21us/step - loss: 0.3475 - acc: 0.8752
Epoch 113/120
58500/58500 [=====] - 1s 21us/step - loss: 0.3468 - acc: 0.8759
Epoch 114/120
58500/58500 [=====] - 1s 21us/step - loss: 0.3462 - acc: 0.8761
Epoch 115/120
58500/58500 [=====] - 1s 21us/step - loss: 0.3453 - acc: 0.8775
Epoch 116/120
58500/58500 [=====] - 1s 19us/step - loss: 0.3444 - acc: 0.8764
Epoch 117/120
58500/58500 [=====] - 1s 19us/step - loss: 0.3438 - acc: 0.8762
Epoch 118/120
58500/58500 [=====] - 1s 21us/step - loss: 0.3431 - acc: 0.8775
Epoch 119/120
58500/58500 [=====] - 1s 16us/step - loss: 0.3420 - acc: 0.8775
Epoch 120/120
58500/58500 [=====] - 1s 19us/step - loss: 0.3418 - acc: 0.8774
```

```
history_dict = history.history
history_dict.keys()
```

```
dict_keys(['loss', 'acc'])
```

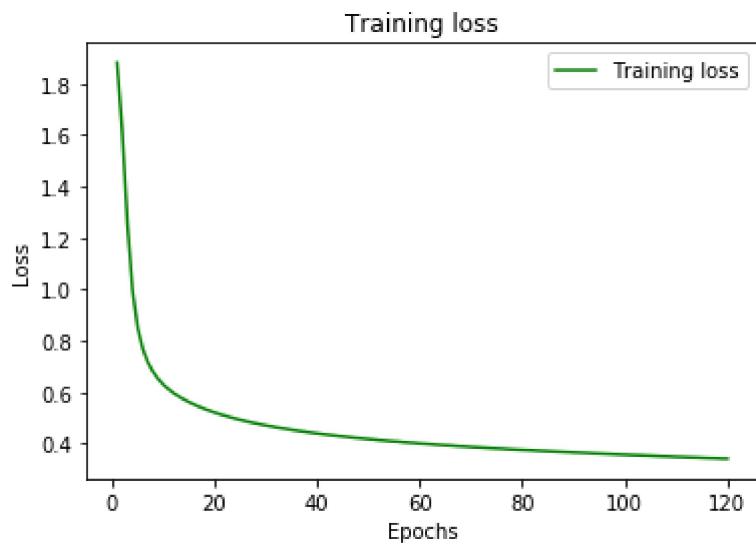
Plot the results

As you might expect, we'll use our `matplotlib` for graphing. Use the data stored in the `history_dict` above to plot the loss vs epochs and the accuracy vs epochs.

```
history_dict = history.history
loss_values = history_dict['loss']

epochs = range(1, len(loss_values) + 1)
plt.plot(epochs, loss_values, 'g', label='Training loss')

plt.title('Training loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.show()
```

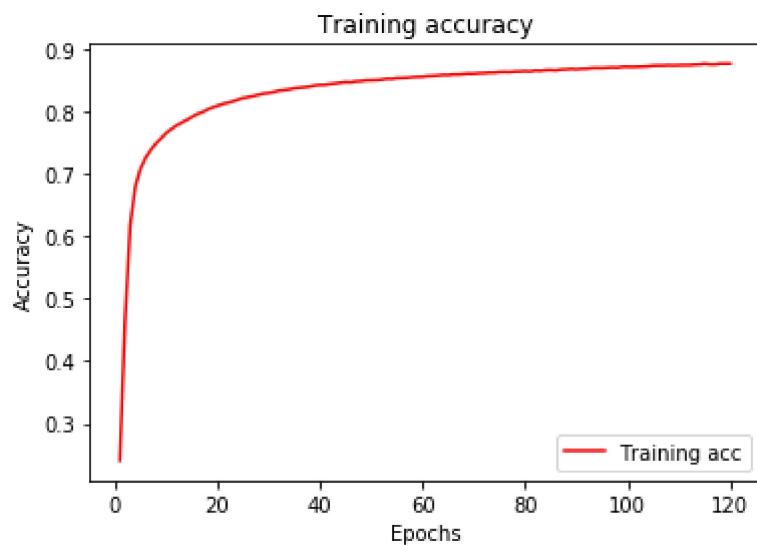


It seems like we could just keep on going and accuracy would go up!

```
# Plot the training accuracy vs the number of epochs

acc_values = history_dict['acc']

plt.plot(epochs, acc_values, 'r', label='Training acc')
plt.title('Training accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



Make predictions

Finally, it's time to make predictions. Use the relevant method discussed in the previous lesson to output (probability) predictions for the test set.

```
# Output (probability) predictions for the test set
y_hat_test = model.predict(test)
```

Evaluate Performance

Finally, print the loss and accuracy for both the train and test sets of the final trained model.

```
# Print the loss and accuracy for the training set
results_train = model.evaluate(train, label_train)
results_train
```

```
58500/58500 [=====] - 2s 29us/step
```

```
[0.33621186860045815, 0.8808376068376068]
```

```
# Print the loss and accuracy for the test set
results_test = model.evaluate(test, label_test)
results_test
```

```
1500/1500 [=====] - 0s 34us/step
```

```
[0.23364664735396704, 0.9266666668256124]
```

We can see that the training set results are really good, and the test set results seem to be even better. In general, this type of result will be rare, as train set results are usually at least a bit better than test set results.

Additional Resources

- https://github.com/susanli2016/Machine-Learning-with-Python/blob/master/Consumer_complaints.ipynb
- <https://catalog.data.gov/dataset/consumer-complaint-database>

Summary

Congratulations! In this lab, you built a neural network thanks to the tools provided by Keras! In upcoming lessons and labs we'll continue to investigate further ideas regarding how to tune and refine these models for increased accuracy and performance.

Releases

No releases published

Packages

No packages published

Contributors 6



Languages

- Jupyter Notebook 100.0%